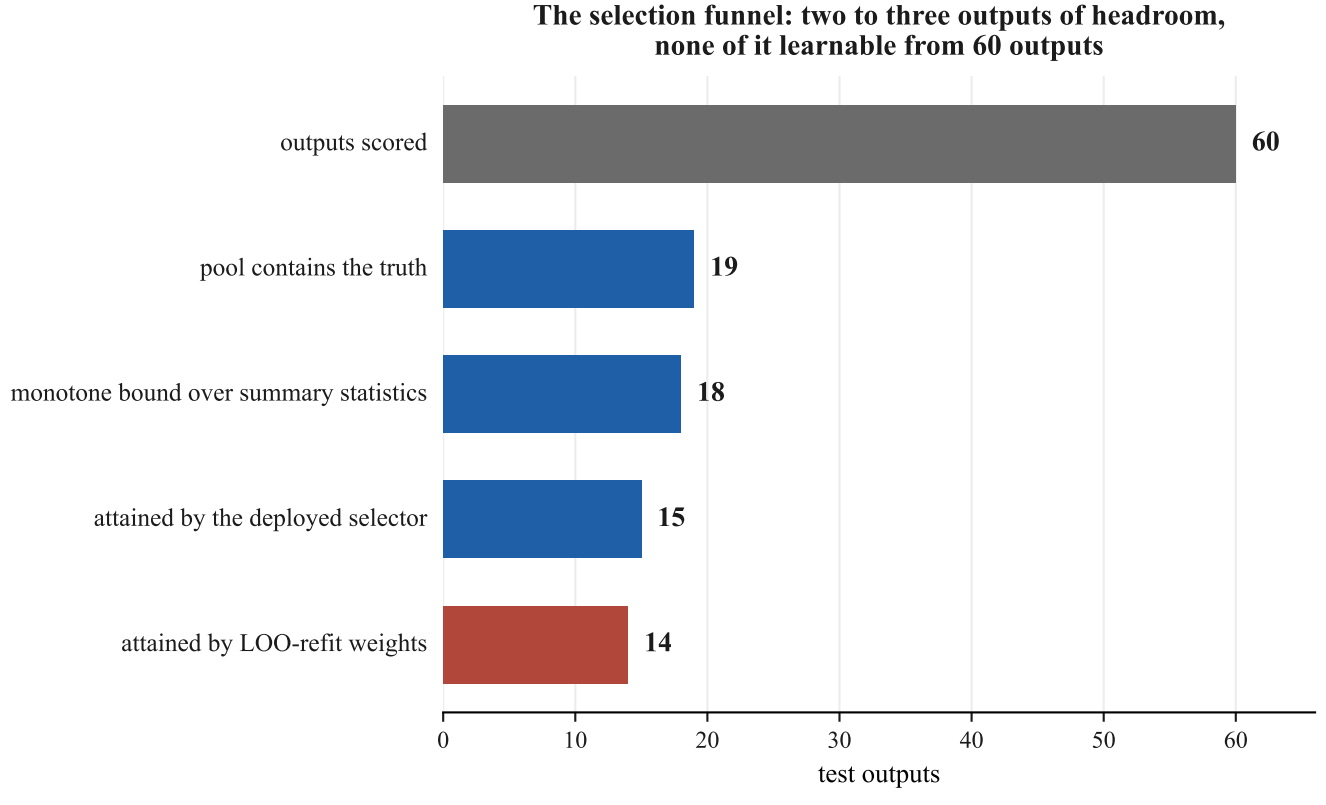
The selection funnel: two to three outputs of headroom, none of it learnable from 60 outputs
outputs scored
pool contains the truth
monotone bound over summary statistics
attained by the deployed selector
attained by LOO-refit weights
60
19
18
15
14
0
10
20
30
40
50
60
test outputs

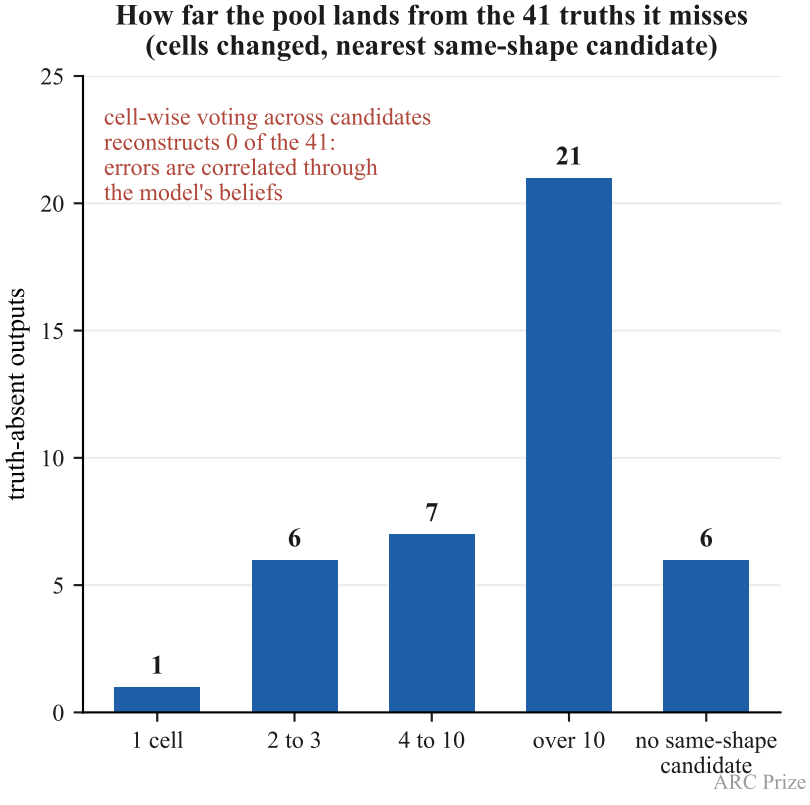
How far the pool lands from the 41 truths it misses (cells changed, nearest same-shape candidate)
cell-wise voting across candidates reconstructs 0 of the 41: errors are correlated through the model's beliefs
truth-absent outputs
0
5
10
15
20
25
1
6
7
21
6
1 cell
2 to 3
4 to 10
over 10
no same-shape candidate